

Mini Project 1

Cedric Lorenz, Leonard Dreessen, Oliver Hess, and Raphael Reimann

Reflection on Manual Query Development

- Without...
 - some experience in query development and
 - access to the schema and additional context

...manually writing the query would not have been possible
- Cycling development of smaller subparts and iteratively testing of those on the database
- Postgres query parser was very helpful while testing
- Manual development felt quite complex, time-consuming and frustrating

Reflection on Development with LLM

- Steps for LLM Generation
 - Provide task and schema to LLM (prompt engineering, < 1 min),
 - Run resulting query (sanity check 2 min)
 - Detailed check of query (10 min)
- LLM-generated query was of high quality
- The LLM was helpful, the result was achieved more quickly and with higher quality
- Follow-up prompts: One initial prompt, quick run of resulting query, then re-prompt to optimize the query which led to faster query execution
- Semi-automatic degree of automation, output and queries had to be manually checked

Generated Query Result (o1-preview)

m_messageid	total_likes	foreign_likes
893353421092	29	28
687194960067	24	17
549756008897	31	21
1030792501923	36	24
687194873730	37	22
893353322965	57	32
962072897868	29	16
1030792344343	32	17
1030792345988	24	12

(9 rows)

Generated Extended Query Result (o1-preview)

m_messageid	total_likes	foreign_likes	first_foreign_liker_name	first_foreign_like_time
893353421092	29	28	Francisco Sanchez	2012-05-09 04:10:33.585+02
687194960067	24	17	Jie Li	2011-10-14 08:43:25.192+02
549756008897	31	21	Eun-Hye Lee	2011-05-29 04:53:09.868+02
1030792501923	36	24	Carlos Parra	2012-09-05 00:33:22.97+02
687194873730	37	22	Baby Yang	2011-10-13 06:24:52.569+02
893353322965	57	32	Miguel Gonzalez	2012-03-21 21:24:21.132+01
962072897868	29	16	Gunnar Johansson	2012-05-09 13:08:46.999+02
1030792344343	32	17	Albert Buysse	2012-08-19 18:19:10.971+02
1030792345988	24	12	Barry Wang	2012-08-05 19:56:31.196+02

(9 rows)

Execution Plans (LLM generated)

LLM-generated

```
Nested Loop Left Join (cost=... rows=... width=...)
Join Filter: (COALESCE(flc.foreign_like_count, 0) >= (ml.total_likes / 2.0))
-> Gather (cost=... rows=... width=...)
Workers Planned: X
-> Hash Join (cost=... rows=... width=...)
Hash Cond: (ml.m_messageid = flc.m_messageid)
-> HashAggregate (cost=... rows=... width=...)
Group Key: ml.m_messageid
-> Nested Loop (cost=... rows=... width=...)
-> HashAggregate (cost=... rows=... width=...)
Group Key: m.m_messageid, m.m_creatorid
Filter: (COUNT(DISTINCT l.l_personid) >= 20)
-> Hash Join (cost=... rows=... width=...)
Hash Cond: (l.l_messageid = m.m_messageid)
-> Seq Scan on likes l (cost=... rows=... width=...)
-> Seq Scan on message m (cost=... rows=... width=...)
Filter: (m_length > 100)
-> Index Scan using likes_pkey on likes l (cost=... rows=1 width=...)
Index Cond: (l_messageid = ml.m_messageid)
-> HashAggregate (cost=... rows=... width=...)
Group Key: flc.m_messageid
-> Hash Join (cost=... rows=... width=...)
Hash Cond: (fl.m_messageid = flc.m_messageid)
-> HashAggregate (cost=... rows=... width=...)
Group Key: fl.m_messageid, fl.liker_id
-> Hash Join (cost=... rows=... width=...)
Hash Cond: (l.m_messageid = fl.m_messageid)
-> Seq Scan on likes l (cost=... rows=... width=...)
-> Seq Scan on foreign_likers fl (cost=... rows=... width=...)
-> Seq Scan on foreign_likes_counts flc (cost=... rows=... width=...)
-> Left Join (cost=... rows=... width=...)
-> Nested Loop Left Join (cost=... rows=... width=...)
-> Nested Loop (cost=... rows=... width=...)
-> Seq Scan on first_foreign_liker flf (cost=... rows=... width=...)
-> Index Scan using person_pkey on person_names pn (cost=... rows=1 width=...)
Index Cond: (p_personid = flf.liker_id)
-> Seq Scan on person_names pn (cost=... rows=... width=...)
```

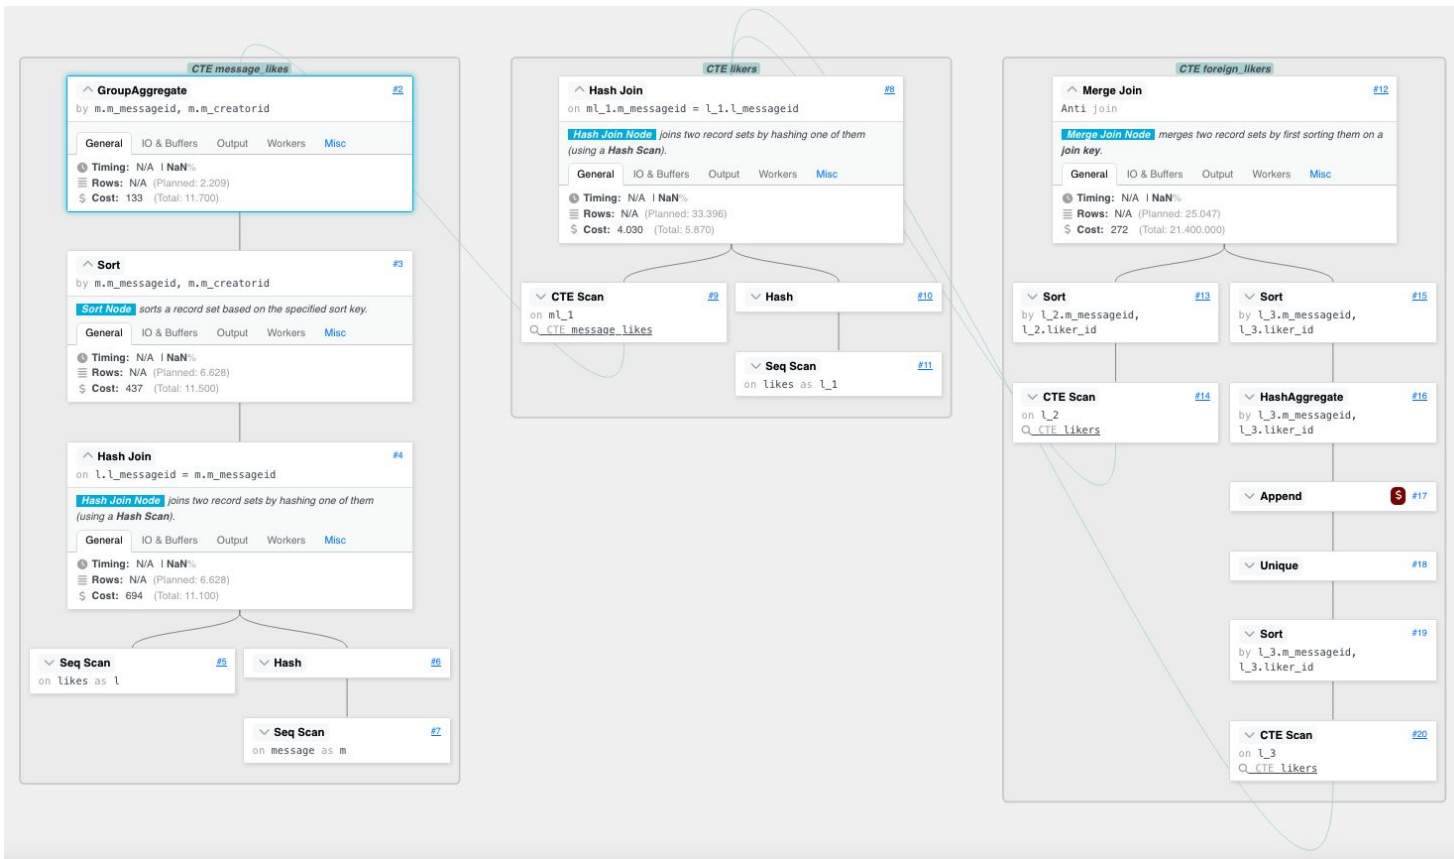
PostgreSQL

```
Sort (cost=21390981.25..21390983.09 rows=736 width=72)
Sort Key: (((COALESCE(flc.foreign_like_count, 0)::bigint))::double precision / (ml.total_likes)::double precision)) DESC
CTE message_likes
-> GroupAggregate (cost=11530.16..11679.29 rows=2209 width=24)
Group Key: m.m_messageid, m.m_creatorid
Filter: (count(DISTINCT l.l_personid) >= 20)
-> Sort (cost=11530.16..11546.73 rows=6628 width=24)
Sort Key: m.m_messageid, m.m_creatorid
-> Hash Join (cost=8840.39..11109.47 rows=6628 width=24)
Hash Cond: (l.l_messageid = m.m_messageid)
-> Seq Scan on likes l (cost=0.00..1792.40 rows=109440 width=16)
-> Hash (cost=8623.30..8623.30 rows=17367 width=16)
-> Seq Scan on message m (cost=0.00..8623.30 rows=17367 width=16)
Filter: (m_length > 100)

CTE likers
-> Hash Join (cost=3802.40..5866.20 rows=33396 width=32)
Hash Cond: (ml_1.m_messageid = l_1.l_messageid)
-> CTE Scan on message_likes ml_1 (cost=0.00..44.18 rows=2209 width=16)
-> Hash (cost=1792.40..1792.40 rows=109440 width=24)
-> Seq Scan on likes l_1 (cost=0.00..1792.40 rows=109440 width=24)

CTE foreign_likers
-> Merge Anti Join (cost=21367216.05..21367591.16 rows=25047 width=24)
Merge Cond: ((l_2.m_messageid = l_3.m_messageid) AND (l_2.liker_id = l_3.liker_id))
-> Sort (cost=3177.19..3260.68 rows=33396 width=24)
Sort Key: l_2.m_messageid, l_2.liker_id
-> CTE Scan on likers l_2 (cost=0.00..667.92 rows=33396 width=24)
-> Sort (cost=21364038.85..21364058.51 rows=7864 width=16)
Sort Key: l_3.m_messageid, l_3.liker_id
-> HashAggregate (cost=21363372.73..21363451.37 rows=7864 width=16)
Group Key: l_3.m_messageid, l_3.liker_id
-> Append (cost=757.58..21363333.41 rows=7864 width=16)
-> Unique (cost=757.58..758.83 rows=163 width=16)
-> Sort (cost=757.58..757.99 rows=167 width=16)
Sort Key: l_3.m_messageid, l_3.liker_id
-> CTE Scan on likers l_3 (cost=0.00..751.41 rows=167 width=16)
Filter: (liker_id = m_creatorid)
```

PostgreSQL execution plan visualized



LLM-generated
could not be
parsed by
explain.dalibo.com

Reflection on the LLM's EXPLAIN output

- LLM faces greater difficulty generating precise, measurement-oriented outputs like PostgreSQL's EXPLAIN statement, compared to the previous tasks
- LLM fails to correctly emulate PostgreSQL's format
- LLM provides broad explanations of identified execution steps, however the output is hardly comprehensible because of lack of visualization